

Forge Neo — Upstream Update Report

2.28.1 → 2.29-17-g76586f6a • 2026-09-10

From	2.28.1
To	2.29-17-g76586f6a
Upstream commits	34
New tags	2.29

Part 1 — Upstream changes

Dependency changes

Package	Was	Now
comfy-kitchen	0.2.31	0.2.33

Commits since last merge (34)

Commit	Subject
76586f6a	Revert "maybe...?"
bd69f173	SolAttn
6bebb747	wrap
efc42fe0	maybe...?
1efa302b	Anima LoRA
84c491cb	oops
4f5f2264	opt
370e4df1	version
1d413a56	vram
8a6bc1d0	lines
a8dcf856	lora-ctl
4bf0bf3b	typing
3f55d7fa	metadata
d2c29a6b	oom
c8147cf1	strip
c2ae52e5	lora-ctl
0d0cb729	arch
0c9273f6	update README
31ea83b1	streamline
3f1702a9	return
92b55e1b	anima
4d3444d6	disable
c5fb32f8	Anima 3.8B
e52cf9f2	dim
28e79335	rewrite LoRA
da635311	shift
9d444214	hires

Commit	Subject
8ce6395e	sort
3206a9ed	iter
a5eb8c58	speed
d70d09ca	speed
d247b40b	ck
22fae2d0	Anima 2.9B
10f7be4e	wipe

Part 2 — Fork merge report

Conflicts resolved

File	Resolution
README.md	Kept ours (Silk Custom readme), per standing rule.
requirements.txt	Took upstream's comfy-kitchen==0.2.33 (newer).
modules/processing.py	Whole-file conflict was a line-ending artifact (our blob was LF, upstream/merge-base CRLF), not real divergence. Diffed with --ignore-cr-at-eol to find upstream's one real change (self.outpath_samples = opts.outdir_hires_samples or self.outpath_samples in the hires-pass sampler) and hand-inserted it into our version, preserving the CUSTOM hr_additional_modules guard, the Wan-model safe-indexing fix, and the 0.60 denoising-strength defaults.
extensions-builtin/sd_forge_lora/networks.py	Real conflict: upstream generalized process_anima() to support the new Anima 2.9B/3.8B block counts (28/40/52) where ours only handled 28/40. Took upstream's version in full (strict superset, not one of our documented customs) and dropped our now-redundant duplicate process_anima() call site. Folder-picker customs (ALL_LORA_FOLDERS, active_lora_dirs, etc.) untouched and verified intact.
backend/attention.py	Clean upstream addition (new @wrap_attn decorator applied to attention_comfy_kitchen_int8, part of a broader attention-module refactor); took upstream's side.
backend/patcher/base.py	Clean upstream addition (wipe_lowvram_weight(m) call in unpatch_model's lowvram branch); took upstream's side.

Custom work verified

Check	Status	Detail
canvas.js custom markers (8/8)	PASS	forgeCanvasInstances, drawWithStamps, _scatterJitterCache, no_shapes, scribbleRotation, scribbleHeight, MAX_HISTORY, _abort all found; file 1988 lines (upstream ~750, untouched by this merge range).
backend/loader.py fp8 TE upcast	PASS	# CUSTOM (Forge Neo) marker present at line 281, untouched by this merge.
modules_forge/utils.py resolution guard	PASS	# CUSTOM (Forge Neo) marker present, untouched by this merge.
modules/generation_parameters_copypaste.py compat shim	PASS	Present, untouched by this merge.
LTX-Video residue	PASS	Precise grep empty - no wiring reintroduced.
Custom ERNIE files	PASS	No ernie_image/ernie_engine files tracked.

Check	Status	Detail
sd-forge-couple tile_funcs.py +3 index shift	PASS	N/A for this merge - that customization was retired 2026-09-04 when sd-forge-couple switched to our Krea2 fork (own repo, unaffected by this forge-neo merge). CLAUDE.md governs as authoritative and documents this as retired, not a failure.

Dependencies installed

- comfy-kitchen==0.2.33

Live test-drive

Launched	yes
Generation	PASS
Tracebacks	0
Notes	Two post-merge bugs found and fixed before this passed (see notes). Krea 2 GGUF checkpoint (sickOllieKrea2GGUF_v10), 512x512, 12 steps, seed 777 via /sdapi/v1/txt2img -> HTTP 200, valid 286KB PNG, 0 tracebacks in tmp/run.log.

Notes

- Two clean auto-merges produced semantically broken (but syntactically valid, non-conflicting) code, caught only by the live launch/generation test:
- 1) backend/args.py: --use-ck-attention was registered twice (once by upstream at its new canonical spot in the attention-flags group, once by our own prior local addition further down in the install-flags group) -> argparse.ArgumentError at launch. Fixed by removing the duplicate, keeping upstream's placement.
- 2) backend/patcher/base.py: our HEAD's unpatch_model() still called reset_weight_functions(m, wipe=True) in a second cleanup loop, but upstream's refactor renamed/replaced that function with wipe_lowvram_weight() (used only inside the lowvram-specific branch) and removed the second loop entirely -> NameError at every model unload/swap (i.e. every generation). Fixed by deleting the now-dead loop to match upstream's current design.
- Both fixes committed as a separate commit after the merge commit, per repo convention (one coherent change per commit).